

JULY 2022

# Model to Predict House Prices : National Case- Shiller Index

**By – Sarth Mirashi**

# Introduction

The primary objective of this study is to build a model to predict the nation-wide US house prices with national Case-Shiller Index as a reference pointer for overall house price trend.

## Publicly available Data Used :

1. **Consumer Price Index CPI** - <https://fred.stlouisfed.org/series/CPIAUCSL>
2. **Unemployment rate** - <https://fred.stlouisfed.org/series/UNRATE>
3. **30-Year Fixed Rate Mortgage** - <https://fred.stlouisfed.org/series/MORTGAGE30US>
4. **Gross Domestic Product** - <https://fred.stlouisfed.org/series/GDP>
5. **Total Unit Labor Cost** - <https://fred.stlouisfed.org/series/LCULMN01USQ661S>
6. **Personal current taxes: State and local: Property taxes-**  
<https://fred.stlouisfed.org/series/S210401A027NBEA>
7. **Employment-Population Ratio** - <https://fred.stlouisfed.org/series/EMRATIO>
8. **NASDAQ Composite Index** - <https://fred.stlouisfed.org/series/NASDAQCOM>
9. **Disposable Personal Income** - <https://fred.stlouisfed.org/series/DSPI>
10. **Rental Vacancy Rate** - <https://fred.stlouisfed.org/series/RRVRUSQ156N>
11. **Monthly Supply of New Houses** - <https://fred.stlouisfed.org/series/MSACSR>
12. **New Privately-Owned Housing Units Completed** -  
<https://fred.stlouisfed.org/series/COMPUTSA>

Code - <https://colab.research.google.com/drive/1fwbMUFUTI8ewy2tdCv0Ro452XS9DcQ8U?usp=sharing>

## DATA:

The above mentioned data was imported and the date was parsed and converted to panda datetime and used as the index of the dataframe. The time-step of the target i.e. Case shiller index is a month thus all the data is preferably collected in this format, if the frequency of the data is higher than monthly, monthly data is used and the rest is discarded. However if the frequency is lower than monthly. Monthly data was filled in using linear interpolation. Quarterly data like GDP were linearly interpolated.


```
factor_quat = {
    'GDP':GDP,
    'Labour_cost':Labour_cost,
    'Rental_vacancy':Rental_vacancy,
    'prop_taxes':prop_taxes,
}
```

```
for name, fact in factor_quat.items(): #interpolation
    df_quat[name].interpolate(method='linear', inplace=True)
```

[illegible]

# Correlation Heatmap :

High inter correlation is seen. This may be attributed to the fact that most of the features considered have a general upward trend with respect to time.



# Regression Approach :

Lagged features expect for the past target variable itself. This will help analyse the target variables dependence on the lagged features (past 1 month ).

## Converting Timeseries forecasting problem to Supervised Learning.:

### Target feature:

Inflation adjusted National Case Shiller index was used as the target variable in the regression approach as it showed better results than using vanilla National CS index.

$$\frac{\text{National Case} - \text{Shiller Index}}{CPI} \times 100$$

### Lagged feature:

All the features used in the regression analysis were lagged features all shift by two months. That is second last month's features were used to predict the target variable. This will allow us to predict two months into the future.


**The previous value of the target variable were not used so as to better study the impact of the factors on the Case shiller index.**

```
target = df1.InfAdj_CSI  
df1 = df1.shift(2)
```

### Train- test split :

The data was split into train and test set to validate the performance. And since this problem is semi - time series problem random sampling for the split is ill advised. This is done to avoid data leakage and overfit.

```
X_train = df1.loc[df1.index < '2012-01-01']  
X_test  = df1.loc[df1.index >= '2012-01-01']  
Y_train = target.loc[df1.index < '2012-01-01']  
Y_test  = target.loc[df1.index >= '2012-01-01']
```



**This split from 2012 stays consistent through all the models.**

## **XGBoost :**

XGboost Algorithm was my go to choice for a semi time series problem. As It is versatile and has seen success in the multivariate time series forecasting.


```
reg = xgb.XGBRegressor(n_estimators=3000, early_stopping_rounds = 50, learning_rate = 0.01)
reg.fit(X_train, Y_train,
        eval_set = [(X_train,Y_train), (X_test,Y_test)],
        verbose = 1000 )
```

[14:57:33] WARNING: /workspace/src/objective/regression\_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

|        |                            |                           |
|--------|----------------------------|---------------------------|
| [0]    | validation_0-rmse:69.8338  | validation_1-rmse:83.5104 |
| [1000] | validation_0-rmse:0.234793 | validation_1-rmse:15.4801 |
| [2000] | validation_0-rmse:0.112942 | validation_1-rmse:15.5299 |
| [2999] | validation_0-rmse:0.067642 | validation_1-rmse:15.4966 |

XGBRegressor(early\_stopping\_rounds=50, learning\_rate=0.01, n\_estimators=3000)

## Feature Importance:




XGBoost allocates feature importance to the features that do well in predicting the target. It was found population is the best in predicting CS index. Considering the feature importance and After iterating through a bunch of feature combinations. Following features were selected in the final model :

|            | Mort_rate | Population | Unemp_rate | house_supply | disp_income | GDP         | Labour_cost | prop_taxes  | month |
|------------|-----------|------------|------------|--------------|-------------|-------------|-------------|-------------|-------|
| date       |           |            |            |              |             |             |             |             |       |
| 1987-03-01 | 9.2040    | 241857.0   | 6.6        | 136.7        | 6159.5      | 4722.156000 | 88.040121   | 2409.000000 | 3     |
| 1987-04-01 | 9.0825    | 242005.0   | 6.6        | 117.8        | 6192.1      | 4750.157333 | 87.713226   | 2420.000000 | 4     |
| 1987-05-01 | 9.0350    | 242166.0   | 6.6        | 126.4        | 6200.0      | 4778.158667 | 87.386331   | 2431.000000 | 5     |
| 1987-06-01 | 9.8325    | 242338.0   | 6.3        | 135.6        | 5967.2      | 4806.160000 | 87.059437   | 2442.000000 | 6     |
| 1987-07-01 | 10.5960   | 242516.0   | 6.3        | 131.4        | 6209.1      | 4832.291667 | 87.039587   | 2459.666667 | 7     |

## Predictions :

The predictions of the model were compared to the test data labelled as true. Root mean squared error of the prediction is **15.5** compared to a standard deviation of 11.5 of the target variable InfAdj\_CSI



## Auto ML:

Pycaret - Is an autoML python library that tests a bunch of different algorithms on the data and lists the best algorithms. It was found that extra trees regressor was best for the purpose.




|          | Model                           | MAE    | MSE          | RMSE   | R2     | RMSLE  | MAPE   | TT (Sec) |
|----------|---------------------------------|--------|--------------|--------|--------|--------|--------|----------|
| et       | Extra Trees Regressor           | 0.3579 | 4.857000e-01 | 0.6255 | 0.9977 | 0.0071 | 0.0044 | 0.456    |
| knn      | K Neighbors Regressor           | 0.4676 | 1.777100e+00 | 0.9981 | 0.9924 | 0.0107 | 0.0056 | 0.062    |
| rf       | Random Forest Regressor         | 0.7098 | 1.756800e+00 | 1.2467 | 0.9910 | 0.0148 | 0.0089 | 0.542    |
| gbr      | Gradient Boosting Regressor     | 0.9729 | 2.377400e+00 | 1.4543 | 0.9882 | 0.0174 | 0.0123 | 0.110    |
| dt       | Decision Tree Regressor         | 0.8447 | 2.954500e+00 | 1.5406 | 0.9840 | 0.0185 | 0.0106 | 0.016    |
| lightgbm | Light Gradient Boosting Machine | 1.1133 | 3.754600e+00 | 1.8152 | 0.9816 | 0.0200 | 0.0134 | 0.070    |
| ada      | AdaBoost Regressor              | 1.8507 | 5.176600e+00 | 2.2668 | 0.9729 | 0.0302 | 0.0255 | 0.108    |
| lr       | Linear Regression               | 3.0086 | 1.691670e+01 | 4.0013 | 0.9144 | 0.0501 | 0.0400 | 0.323    |

## Extra Trees Regressor :

As suggest by Pycaret Extra trees regressor was applied to the problem.


## Feature importance:



The feature importance observed in Extra trees was different from XGBoost

## Predictions :

The predictions of the model were compared to the test data labelled as true. Root mean squared error of the prediction is **13.9** compared to a standard deviation of 11.5 of the target variable InfAdj\_CSI



## Bagging Ensemble :

In this technique Trained models are assigned weights according to their performance and weighted average of their predictions is considered as the prediction of the ensemble. The models used in the ensemble are as follows :

```
✓ [57] models = {'gbr': GradientBoostingRegressor(n_estimators=10000, learning_rate = 0.01),  
0s      'br': BayesianRidge(),  
      'xgb': xgb.XGBRegressor(n_estimators=3000, early_stopping_rounds = 50, learning_rate = 0.001, verbose = 0),  
      'etr': ExtraTreesRegressor(n_estimators=3000)}  
  
✓ [58] for name, model in models.items():  
11s      model.fit(X_train, Y_train)
```

### Cross validation:

KFold cross validation was used. The root mean square value for each individual model is as mentioned in the figure


```
results = {}  
  
kf = KFold(n_splits=10)  
  
for name, model in models.items():  
    result = np.sqrt(-cross_val_score(model, df1, target, scoring='neg_mean_squared_error', cv=kf))  
    results[name] = result
```

```
-----  
gbr  
5.190255777086366  
4.482449281626742  
-----  
br  
7.346395731053842  
3.8106547123155434  
-----  
xgb  
7.452391132553402  
6.537198307354843  
-----  
etr  
4.208012381326437  
4.504610177953215
```

```
y_pred = (  
    0.15 * models['gbr'].predict(X_test) +  
    0.25 * models['br'].predict(X_test) +  
    0.4 * models['etr'].predict(X_test) +  
    0.2 * models['xgb'].predict(X_test))
```

## Predictions :

The predictions of the model were compared to the test data labelled as true. Root mean squared error of the prediction is **4.64** compared to a standard deviation of 11.5 of the target variable InfAdj\_CSI



# Vector autoregression:

Vector Autoregression (VAR) is a forecasting algorithm that can be used when two or more time series influence each other. That is, the relationship between the time series involved is bi-directional

## Darts Library :

Darts is a Python library for easy manipulation and forecasting of time series. It contains a variety of models, from classics such as ARIMA to deep neural networks.

**ADfuller test :** Adfuller test is used to check for stationarity of the data. The ADfuller test was performed .The data as is not stationary however , after differencing the data once. The data is not stationary

Differencing :

```
def adftest(df):  
  
    for fact in df:  
        print('\n-----{}-----\n'.format(fact))  
        adf = adfuller(df[fact])  
        print(f'ADF Statistic: {adf[0]}')  
        print(f'p-value: {adf[1]}')  
  
    adftest(df2.diff(1)[1:])
```

ADfuller test results :

```
-----InfAdj_CSI-----  
ADF Statistic: -2.9127271274156596  
p-value: 0.043887896852186636  
  
-----Mort_rate-----  
ADF Statistic: -10.322781972846803  
p-value: 2.9788996059807623e-18  
  
-----Unemp_rate-----  
ADF Statistic: -12.1275063492022  
p-value: 1.7667530500682174e-22  
  
-----disp_income-----  
ADF Statistic: -4.462838431231033  
p-value: 0.0022938490894836274  
  
-----GDP-----  
ADF Statistic: -8.749766128154748  
p-value: 2.857550554251797e-14  
  
-----prop_taxes-----  
ADF Statistic: -3.7979600886659948  
p-value: 0.002928255045692676
```

**Thus stationary**

## Varima Model :

After some hyperparameter tuning following parameters were selected for the final model :


$P = 3$ ,  $d = 1$ ,  $q = 0$ . Since  $q$  was zero the moving average component was not used making the model a VAR model.

```
model1 = VARIMA(p = 4, d = 1, q = 0)
model1.fit(X_train)
pred1 = model1.predict(len(X_test))
pred1.to_csv('varima')
final = pd.read_csv('/content/varima')
final.index = pd.to_datetime(final['date'])
```

- **p** (*int*) – Order (number of time lags) of the autoregressive model (AR)
- **d** (*int*) – The order of differentiation; i.e., the number of times the data have had past values subtracted. (I) Note that Darts only supports  $d \leq 1$  because for  $d > 1$  the optimizer often does not result in stable predictions. If results are not stable for  $d = 1$  try to set  $d = 0$  and enable the trend parameter to account for possible non-stationarity.
- **q** (*int*) – The size of the moving average window (MA).

## Predictions :

The predictions of the model were compared to the test data labelled as true. In this case predictions were made about **10 years in the future**. Root mean squared error of the prediction is 9.06 compared to a standard deviation of 11.5 of the target variable InfAdj\_CSI



# LSTM :

LSTM model of the following specifications was developed however failed to delivered due to some errors.

Model: "sequential\_2"

| Layer (type)              | Output Shape   | Param # |
|---------------------------|----------------|---------|
| lstm_6 (LSTM)             | (None, 4, 128) | 70144   |
| leaky_re_lu_4 (LeakyReLU) | (None, 4, 128) | 0       |
| lstm_7 (LSTM)             | (None, 4, 128) | 131584  |
| leaky_re_lu_5 (LeakyReLU) | (None, 4, 128) | 0       |
| dropout_4 (Dropout)       | (None, 4, 128) | 0       |
| lstm_8 (LSTM)             | (None, 64)     | 49408   |
| dropout_5 (Dropout)       | (None, 64)     | 0       |
| dense_2 (Dense)           | (None, 1)      | 65      |
| Total params: 251,201     |                |         |
| Trainable params: 251,201 |                |         |
| Non-trainable params: 0   |                |         |



